

# REINFORCEMENT LEARNING

Machine Learning  
Illy CSE IDP and IDDMP

Refer : Stephen Marsland

# Motivation

- **Gap in Learning Paradigms**

- Supervised learning: trained on correct answers
- Unsupervised learning: exploits only similarities to cluster
- Reinforcement learning sits in between, receiving only success/failure feedback

- **Need for Trial-and-Error**

- Learner must try different strategies without explicit improvement instructions
- “Trying out” strategies is equivalent to performing a search over state/action space

- **Role of Search**

- Search underlies the process of testing strategies to maximize a reward
- Connects back to general search methods

# What Is Reinforcement Learning?

## Agent–Environment Interaction

1. **Agent:** the learner making decisions
2. **Environment:** the world providing states and rewards

## Reward Function

3. Signals “how good” a chosen strategy is
4. Guides the learner to repeat satisfying actions and discard others

## Trial-and-Error & the Law of Effect

5. Actions followed by satisfaction become more likely; those followed by discomfort become less likely
6. Originates from Thorndike (1911): strengthening or weakening of action–situation bonds based on outcome

## Delayed Credit Assignment

7. Not always the last action that causes an outcome (e.g. child waving arms before falling)
8. Challenges in determining which earlier actions led to success or failure

# Reinforcement Learning Framework

## Objective

Map states (inputs) to actions so as to maximize a numerical reward

## Agent vs. Environment

**Agent:** selects actions based on current state

**Environment:** produces states (sensor readings) and rewards

## Robot Example

**State:** sensor readings (noisy, partial view of world)

**Actions:** motor commands that move the robot

**Reward:** task performance (e.g. traversing without crashing)

## Delayed Rewards

Reward may arrive long after the causative actions (e.g. reaching maze center)

Must consider both immediate and future expected rewards

## Policy & Decision Trade-off

**Policy:** mapping from state  $\rightarrow$  action

Balance **exploitation** (best known action) vs. **exploration** (trying new actions)

# Reinforcement Learning Framework (Robot Example)




FIGURE 11.1 A robot perceives the current state of its environment through its sensors, and performs actions by moving its motors. The reinforcement learner (agent) within the robot tries to predict the next state and reward.

# Reinforcement Learning Cycle

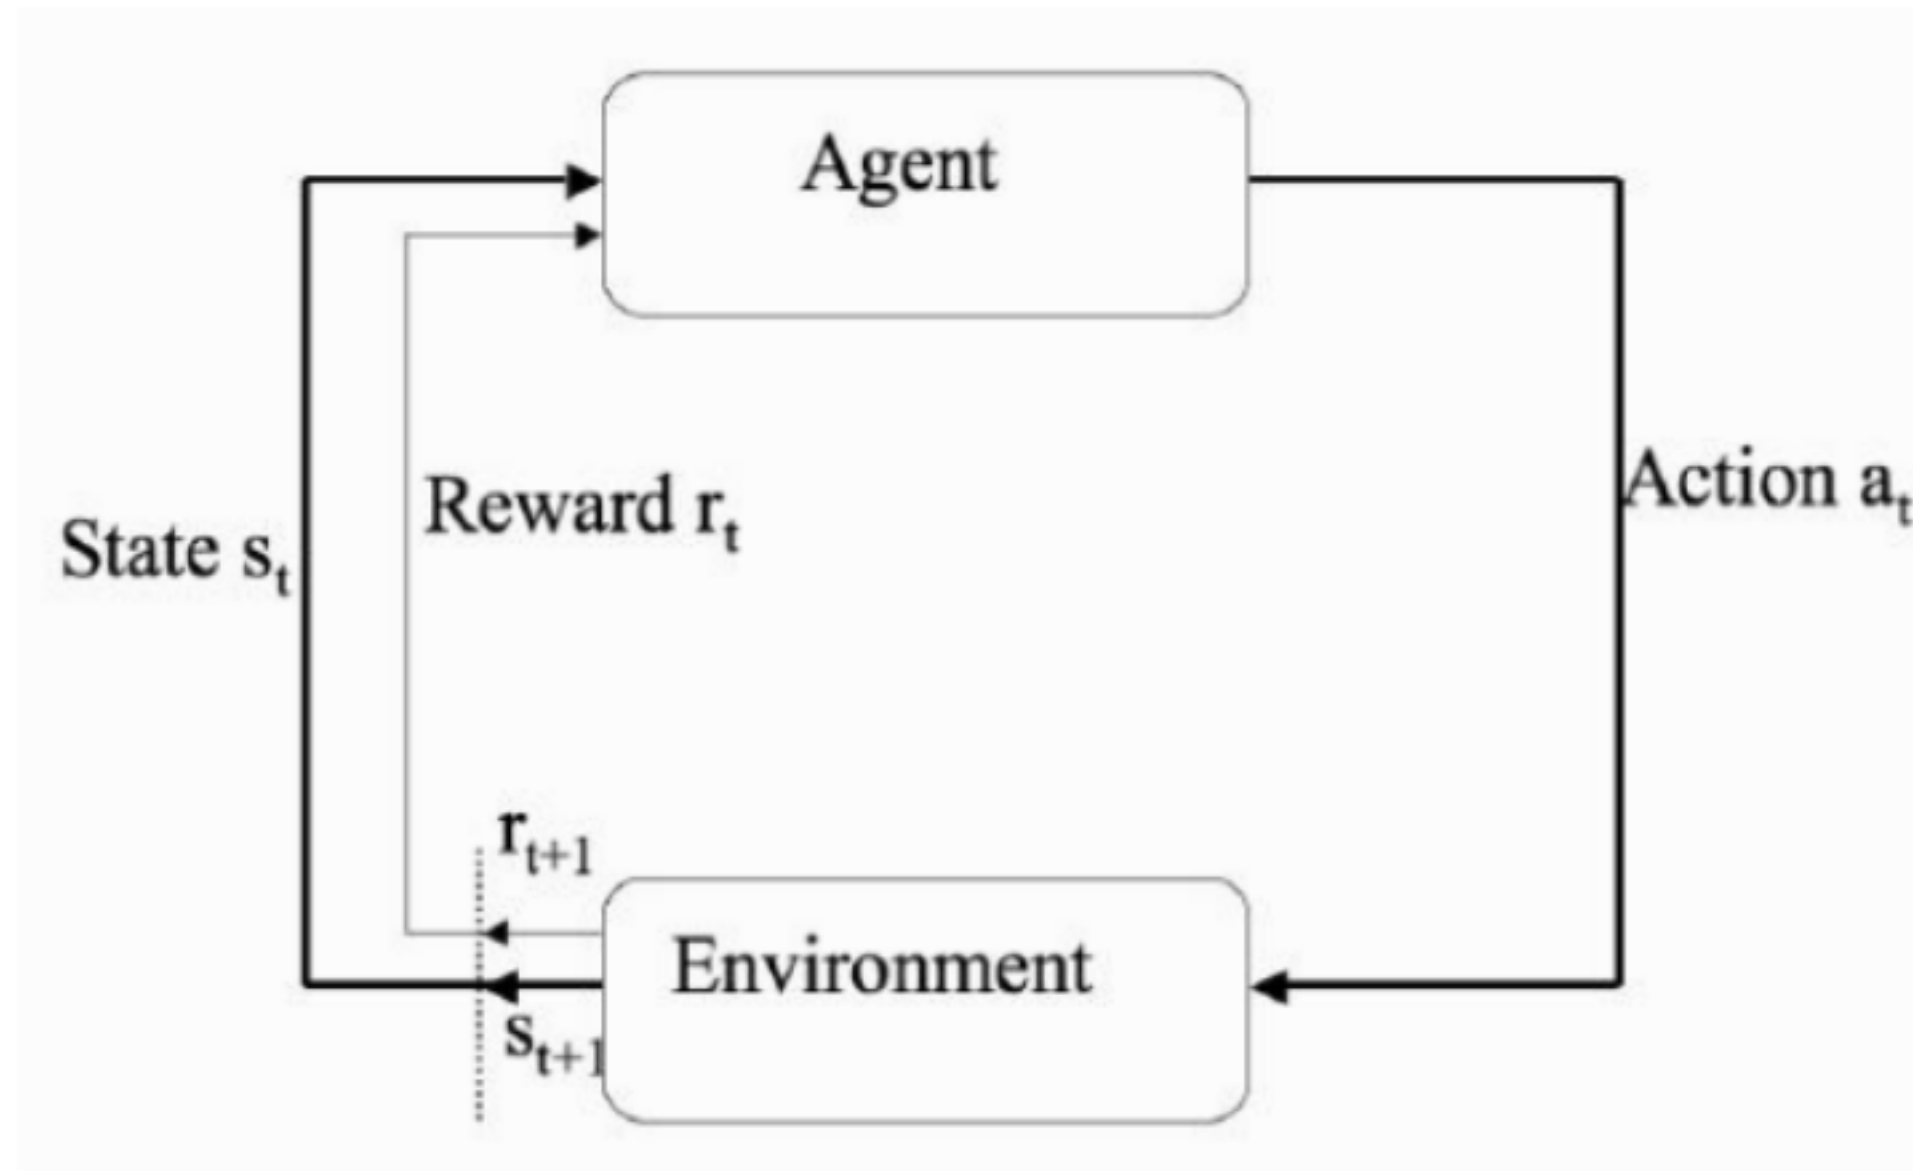


FIGURE 11.2 The reinforcement learning cycle: the learning agent performs action  $a_t$  in state  $s_t$  and receives reward  $r_{t+1}$  from the environment, ending up in state  $s_{t+1}$ .

# Example – Getting Lost

## Scenario

1. Arrive in unfamiliar old-town district at 3 a.m., completely lost
2. Draw map of connected squares (states A–F)
3. Goal: find hostel in square F

## Reward Design

4. **Delayed reward:** only upon reaching F (eat all chips)
5. **Stay still** (sleep on feet): reward =  $-5$  (punishment to encourage movement)
6. **All other moves:** reward = 0 (neutral)
7. **Invalid moves** (no direct road): no action possible

## Absorbing State

State F is absorbing (once entered, stay and receive final reward)

# Old Town Graph




FIGURE 11.3 The old town that you find yourself lost in.



# State Diagram of Old Town Graph




FIGURE 11.4 The *state diagram* if you are correct and the backpacker's is in square (state) F. The connections from each state back into itself (meaning that you don't move) are not shown, to avoid the figure getting too complicated. They are each worth  $-5$  (except for staying in state F, which means that you are in the backpacker's).

# Reward Matrix

Current State	Next State					
	A	B	C	D	E	F
A	-5	0	-	-	-	-
B	0	-5	0	0	-	-
C	-	0	-5	0	-	100
D	-	0	0	-5	0	-
E	-	-	-	0	-5	100
F	-	-	0	-	0	-

# State & Action Spaces

**State Space:** all possible states the agent can experience

Size grows with number of inputs and their ranges

Example: 5 inputs each 0 – 100  $\rightarrow 100^5$  states (huge)

**Action Space:** all possible actions available in each state

**Dimensionality Reduction:**

Quantization can shrink state space

Example: map each input to two classes ( $<50, \geq 50$ )  $\rightarrow 2^5 = 32$  states

**Trade-Off:**

Smaller spaces speed up learning

Oversimplification may lose important information

**Key Insight:** carefully choose and balance state/action spaces to make learning tractable without sacrificing accuracy.

# Carrots & Sticks – The Reward Function

**Purpose:** guide the learner by assigning numerical rewards to (state, action) pairs

**Design:** akin to a fitness function—crafted ad hoc to reflect desired behavior

**External Signal:** rewards come from the environment, not the learner

**Positive & Negative:**

- Positive reward → reinforcement of action

- Negative reward (“punishment”) → actions to avoid

**Avoid Sub-Goals:** extra incentives may be exploited without achieving the real objective

**Reward Variants (Maze Example):**

- +50 at goal only → learn to reach center

- 1 per move +50 at goal → bias toward shorter paths

**Task Types:**

- Episodic:** finite episodes, propagate end reward backward

# Action Selection

## Estimating Action Value (Q)

$Q_{s,t}(a)$  : average reward observed when taking action  $a$  in state  $s$  over  $t$  trials

Converges to the true expected reward with enough samples

## Greedy

Always pick the action with the highest  $Q_{s,t}(a)$

Pure exploitation, no exploration

## $\epsilon$ -Greedy

With probability  $1-\epsilon$  : pick the greedy action

With probability  $\epsilon$ : pick a random action  $\rightarrow$  injects exploration

## Soft-Max

Assigns selection probabilities proportional to  $\exp(Q_{s,t}(a)/\tau)$

Temperature  $\tau$  controls exploration vs. exploitation

Large  $\tau$ : nearly uniform choice

Small  $\tau$ : focuses on higher-valued actions

$$P(Q_{s,t}(a)) = \frac{\exp(Q_{s,t}(a)/\tau)}{\sum_b \exp(Q_{s,t}(b)/\tau)}.$$

# Action Selection

**Definition:** a policy  $\pi$  is a (possibly deterministic) mapping from each state  $s$  to an action  $a$

**Purpose:** defines the agent's behavior by specifying which action to take in every state

**Optimal Policy:** one that maximizes expected future (discounted) reward

**Learning the Policy:** core challenge of reinforcement learning

**Key Considerations:**

**History Dependence:** how much past information ( $st_0, st_1, \dots$ ) is needed to choose the best action now (Markov property)

**State Valuation:** how to assign a value to the current state to guide policy improvement



# Summary & Conclusion

## Reinforcement Learning Essence:

Learner (agent) interacts with an environment, mapping states to actions to maximize rewards

## Core Components:

**State & Action Spaces:** define what the agent perceives and can do

**Reward Function:** external signal guiding desired behavior

**Policy:** strategy mapping states to actions

## Trade-Offs:

**Exploration vs. Exploitation:** balance trying new actions and using known good ones

**Discounting:** prioritizing near-term rewards when future outcomes are uncertain

## Design Challenges:

Crafting appropriate reward functions without unintended shortcuts

Managing large state/action spaces through abstraction or quantization

## Goal:

Learn an optimal policy that yields the highest cumulative reward over time, adapting through trial and error.